

\*\*\*\*\*

## *Cassandra :: Multi-Node Cluster*

\*\*\*\*\*

[Apache Cassandra](#) is a highly scalable open source database system, achieving great performance on multi-node setups.

Previously, we went over [how to run a single-node Cassandra cluster](#). In this tutorial, you'll learn how to install and use Cassandra to run a multi-node cluster on Ubuntu 14.04.

### *Prerequisites*

1) *ssh user without password authentication*

2) *install cassandra on all nodes with single-node installation*

in my case, i used 2 nodes.

1) 192.168.1.253      *ubuntu-server.com*    *server*  
2) 192.168.1.254      *ubuntu-client.com*    *client*  
*ssh-user without password authentication*    *ubuntu*

\*\*\*\*\*

## *Setup SSH for Auto Login without a Password*

\*\*\*\*\*

on both server :

*# adduser ubuntu*

give all details for this user in both machine

*# usermod -aG sudo ubuntu*

*# visudo*

add this line at the bottom of this file

*ubuntu ALL=(ALL) NOPASSWD: ALL*

*# /etc/init.d/ssh restart*

*# su ubuntu*

*# cd /home/ubuntu/.ssh/*

*# ssh-keygen -t rsa*

```
ubuntu@akhil: ~/.ssh
ubuntu@akhil: ~
ubuntu@akhil:~/.ssh$ ls
known_hosts
ubuntu@akhil:~/.ssh$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/ubuntu/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/ubuntu/.ssh/id_rsa.
Your public key has been saved in /home/ubuntu/.ssh/id_rsa.pub.
The key fingerprint is:
c4:97:16:02:bd:73:8c:db:cb:58:aa:75:b5:9c:29:23 ubuntu@akhil.bizruntime.com
The key's randomart image is:
+--[ RSA 2048 ]-----+
|.O. .|
|... o|
|o++|
|.oo|
|S= .|
|.oo +|
|E+=.=|
|.oooo|
|..|
+-----+
ubuntu@akhil:~/.ssh$
```

on server-side:

```
# sudo ssh-copy-id 192.168.1.254
```

on client-side:

```
# cd /home/ubuntu/.ssh/
# cat id_rsa.pub >> authorized_keys
```

copy this output

on server-side:

```
# cd /home/ubuntu/.ssh/
# nano authorized_keys
```

paste that in this file.

On both server :

```
# sudo /etc/init.d/ssh restart
```

```
*****
links: http://www.rebol.com/docs/ssh-auto-login.html
```

```
*****
```

\*\*\*\*\*

## *Cassandra Configuration*

\*\*\*\*\*

On both machine do this,

```
# su ubuntu
```

Cassandra requires that the Oracle Java SE Runtime Environment (JRE) be installed. So, in this step, you'll install and verify that it's the default JRE.

```
# sudo add-apt-repository ppa:webupd8team/java
```

```
# sudo apt-get update
```

Then install the Oracle JRE. Installing this particular package not only installs it but also makes it the default JRE. When prompted, accept the license agreement:

```
# sudo apt-get install oracle-java8-set-default
```

```
# java -version
```

You should see output similar to the following:

### **Output**

```
java version "1.8.0_60"
```

```
Java(TM) SE Runtime Environment (build 1.8.0_60-b27)
```

```
Java HotSpot(TM) 64-Bit Server VM (build 25.60-b23, mixed mode)
```

We'll install Cassandra using packages from the official Apache Software Foundation repositories, so start by adding the repo so that the packages are available to your system. Note that Cassandra 2.2.2 is the latest version at the time of this publication. Change the 22x to match the latest version. For example, use 39x if Cassandra 3.9 is the latest version:

```
# echo "deb http://www.apache.org/dist/cassandra/debian 39x main" | sudo tee -a /etc/apt/sources.list.d/cassandra.sources.list
```

```
# echo "deb-src http://www.apache.org/dist/cassandra/debian 39x main" | sudo tee -a /etc/apt/sources.list.d/cassandra.sources.list
```

\*\*\*\*\*

**links ::**

**latest version ::**

<http://cassandra.apache.org/download/>

\*\*\*\*\*

```
# sudo apt-get update
```

Finally, install Cassandra:

```
# sudo apt-get install cassandra
```

Ordinarily, Cassandra should have been started automatically at this point. However, because of a bug, it does not. To confirm that it's not running, type:

```
# sudo service cassandra status  
# sudo service cassandra restart
```

wait for few minutes and check,

```
# sudo tailf /var/log/cassandra/system.log
```

```
# sudo service cassandra status
```

If it is not running, the following output will be displayed:

Output

```
* could not access pidfile for Cassandra
```

and check the log file also which mentioned above.

Then, resolve this issue by,

```
# sudo nano +60 /etc/init.d/cassandra
```

That line should read:

```
# sudo nano /etc/init.d/cassandra
```

```
CMD_PATT="cassandra.+CassandraDaemon"
```

Change it to:

```
CMD_PATT="cassandra"
```

Close and save the file, then reboot the server:

```
# sudo reboot
```

After logging back in, Cassandra should now be running. Verify:

```
# sudo service cassandra status
```

If you are successful, you will see:

Output

*\* Cassandra is running*

If you were able to successfully start Cassandra, check the status of the cluster:

*# sudo nodetool status*

In the output, **UN** means it's Up and Normal:

Output

***Datacenter: datacenter1***

*=====*

***Status=Up/Down***

***/ State=Normal/Leaving/Joining/Moving***

<i>--</i>	<i>Address</i>	<i>Load</i>	<i>Tokens</i>	<i>Owns</i>	<i>Host ID</i>	<i>Rack</i>
<i>UN</i>	<i>127.0.0.1</i>	<i>142.02 KB</i>	<i>256</i>	<i>?</i>	<i>2053956d-7461-41e6-8dd2-0af59436f736</i>	<i>rack1</i>

***Note: Non-system keyspaces don't have the same replication settings, effective ownership information is meaningless***

Then connect to it using its interactive command line interface `cqlsh`.

*# cqlsh*

You will see it connect:

Output

***Connected to Test Cluster at 127.0.0.1:9042.***

***[cqlsh 5.0.1 | Cassandra 2.2.2 | CQL spec 3.3.1 | Native protocol v4]***

***Use HELP for help.***

***cqlsh>***

Type `exit` to quit:

*> exit*

*\*\*\*\*\**

***links ::***

*<https://www.digitalocean.com/community/tutorials/how-to-install-cassandra-and-run-a-single-node-cluster-on-ubuntu-14-04>*

*\*\*\*\*\**

The first command you'll run on each node will stop the Cassandra daemon.

```
# sudo service cassandra stop
```

When that's completed, delete the default dataset.

```
# sudo rm -rf /var/lib/cassandra/data/*  
# sudo mkdir /var/lib/cassandra/data/system  
# sudo chown -R cassandra:cassandra /var/lib/cassandra/data
```

## *Configuring the Cluster*

Cassandra's configuration file is located in the `/etc/cassandra` directory. That configuration file, `cassandra.yaml`, contains many directives and is very well commented. In this step, we'll modify that file to set up the cluster.

Only the following directives need to be modified to set up a multi-node Cassandra cluster:

- `cluster_name`: This is the name of your cluster.
- `-seeds`: This is a comma-delimited list of the IP address of each node in the cluster.
- `listen_address`: This is IP address that other nodes in the cluster will use to connect to this one. It defaults to **localhost** and needs changed to the IP address of the node.
- `rpc_address`: This is the IP address for remote procedure calls. It defaults to **localhost**. If the server's hostname is properly configured, leave this as is. Otherwise, change to server's IP address or the loopback address (127.0.0.1).
- `endpoint_snitch`: Name of the snitch, which is what tells Cassandra about what its network looks like. This defaults to **SimpleSnitch**, which is used for networks in one datacenter. In our case, we'll change it to **GossipingPropertyFileSnitch**, which is preferred for production setups.
- `auto_bootstrap`: This directive is not in the configuration file, so it has to be added and set to **false**. This makes new nodes automatically use the right data. It is optional if you're adding nodes to an existing cluster, but required when you're initializing a fresh cluster, that is, one with no data.

Open the configuration file for editing using nano or your favorite text editor.

```
# sudo nano /etc/cassandra/cassandra.yaml
```

Search the file for the following directives and modify them as below to match your cluster. Replace `your_server_ip` with the IP address of the server you're currently working on. The `- seeds:` list should be the same on every server, and will contain each server's IP address separated by commas.

...

*`cluster_name: 'CassandraDOCluster'`*

...

*`seed_provider:`*

*`- class_name: org.apache.cassandra.locator.SimpleSeedProvider`*

*`parameters:`*

*`- seeds: "your_server_ip,your_server_ip_2,...your_server_ip_n"`*

...

*`listen_address: your_server_ip`*

...

*`rpc_address: your_server_ip`*

...

*`endpoint_snitch: GossipingPropertyFileSnitch`*

...

*At the bottom of the file, add in the `auto_bootstrap` directive by pasting in this line:*

*`auto_bootstrap: false`*

*example :*

*on server-side:*

*...*

*cluster\_name: 'bizruntime'*

*...*

*seed\_provider:*

*- class\_name: org.apache.cassandra.locator.SimpleSeedProvider*

*parameters:*

*- seeds: "192.168.1.253"*

*...*

*listen\_address: 192.168.1.253*

*...*

*rpc\_address: 192.168.1.253*

*...*

*endpoint\_snitch: GossipingPropertyFileSnitch*

*...*

*At the bottom of the file, add in the auto\_bootstrap directive by pasting in this line:*

*auto\_bootstrap: false*



*on client -side:*

...

*cluster\_name: 'bizruntime'*

...

*seed\_provider:*

*- class\_name: org.apache.cassandra.locator.SimpleSeedProvider*

*parameters:*

*- seeds: "192.168.1.253"*

...

*listen\_address: 192.168.1.254*

...

*rpc\_address: 192.168.1.254*

...

*endpoint\_snitch: GossipingPropertyFileSnitch*

...

At the bottom of the file, add in the auto\_bootstrap directive by pasting in this line:

auto\_bootstrap: false

*# sudo service cassandra start*

If you check the status of the cluster, you'll find that only the local node is listed, because it's not yet able to communicate with the other nodes.

*# sudo nodetool status*

*Datacenter: dc1*

*=====*

*Status=Up/Down*

*|/ State=Normal/Leaving/Joining/Moving*

<i>--</i>	<i>Address</i>	<i>Load</i>	<i>Tokens</i>	<i>Owns</i>	<i>Host ID</i>	<i>Rack</i>
<i>UN</i>	<i>192.168.1.253</i>	<i>147.48 KB</i>	<i>256</i>	<i>?</i>	<i>f50799ee-8589-4eb8-a0c8-241cd254e424</i>	<i>rack1</i>
<i>UN</i>	<i>192.168.1.254</i>	<i>139.04 KB</i>	<i>256</i>	<i>?</i>	<i>54b16af1-ad0a-4288-b34e-cacab39caeec</i>	<i>rack1</i>

*Note: Non-system keyspaces don't have the same replication settings, effective ownership information is meaningless*

You can also check if you can connect to the cluster using cqlsh, the Cassandra command line client. Note that you can specify the IP address of any node in the cluster for this command.

```
# cqlsh your_server_ip 9042
```

example :

```
# cqlsh 192.168.1.253 9042
```

```
Connected to bizruntime at 192.168.1.253:9042.
[cqlsh 5.0.1 | Cassandra 3.9 | CQL spec 3.4.2 | Native protocol v4]
Use HELP for help.
cqlsh>
```

```
*****
links ::
https://www.digitalocean.com/community/tutorials/how-to-run-a-multi-node-cluster-database-with-cassandra-on-ubuntu-14-04
*****
```

```
*****
testing ::
*****
```

login

```
# cqlsh 192.168.1.253 9042
```

create keyspace with replication

```
> CREATE KEYSPACE biz WITH REPLICATION = { 'class' : 'SimpleStrategy',
'replication_factor' : 2 };
```

check on both server

```
> SELECT * FROM system_schema.keyspaces;
```

to use that keyspace (database)

```
> use biz;
```

create tables

```
> use biz;
> create table emp (empid int primary key,
... emp_first varchar, emp_last varchar, emp_dept varchar);
```

add another values to that table

```
> insert into emp (empid, emp_first, emp_last, emp_dept)
... values (1,'fred','smith','eng');
```

check from both server

```
> select * from emp;
```

```
empid | emp_dept | emp_first | emp_last
-----+-----+-----+-----
1 | eng | fred | smith
```

(1 rows)

\*\*\*\*\*

link ::

to create tables >>

<http://datascale.io/cassandra-partitioning-and-clustering-keys-explained/>

\*\*\*\*\*

\*\*\*\*\*

## *multiple datacenters*

\*\*\*\*\*

This topic contains information for deploying a Cassandra cluster with multiple data centers. In Cassandra, the term datacenter is a grouping of nodes. datacenter is synonymous with replication group, that is, a grouping of nodes configured together for replication purposes.

### *Prerequisites*

Each node must be correctly configured before starting the cluster. You must determine or perform the following before starting the cluster:

- A good understanding of how Cassandra works. Be sure to read at least [Understanding the architecture](#), [Data replication](#), and [Cassandra's rack feature](#).
- Install Cassandra on each node.
- Choose a name for the cluster.
- Get the IP address of each node.
- Determine which nodes will be seed nodes. **Do not make all nodes seed nodes.** Please read [Internode communications \(gossip\)](#).
- Determine the [snitch](#) and [replication strategy](#). The [GossipingPropertyFileSnitch](#) and [NetworkTopologyStrategy](#) are recommended for production environments.
- If using multiple datacenters, determine a naming convention for each data center and rack, for example: DC1, DC2 or 100, 200 and RAC1, RAC2 or R101, R102. Choose the name carefully; renaming a datacenter is not possible.
- Other possible configuration settings are described in [cassandra.yaml configuration file](#) and property files such as `cassandra-rackdc.properties`.

## Procedure

1. Suppose you install Cassandra on these nodes:

**node0 10.168.66.41 (seed1)**

**node1 10.176.43.66**

**node2 10.168.247.41**

**node3 10.176.170.59 (seed2)**

**node4 10.169.61.170**

**node5 10.169.30.138**

Note: It is a best practice to have more than one seed node per datacenter.

2. If you have a firewall running in your cluster, you must open certain ports for communication between the nodes. See [Configuring firewall port access](#).
3. If Cassandra is running, you must stop the server and clear the data:

Doing this removes the default [cluster\\_name](#) (Test Cluster) from the system table. All nodes must use the same cluster name.

Package installations:

- a. Stop Cassandra:

**# *sudo service cassandra stop***

- b. Clear the data:

**# *sudo rm -rf /var/lib/cassandra/data/system/\****

4. Set the properties in the [cassandra.yaml](#) file for each node:

Note: After making any changes in the `cassandra.yaml` file, you must restart the node for the changes to take effect.

Properties to set:

- a. `num_tokens`: *recommended value: 256*

- b. `-seeds`: *internal IP address of each seed node*

Seed nodes do not [bootstrap](#), which is the process of a new node joining an existing cluster. For new clusters, the bootstrap process on seed nodes is skipped.

- c. `listen_address`:

If not set, Cassandra asks the system for the local address, the one associated with its hostname. In some cases Cassandra doesn't produce the correct address and you must specify the `listen_address`.

- d. `endpoint_snitch`: *name of snitch* (See [endpoint snitch](#).) If you are changing snitches, see [Switching snitches](#).

- e. `auto_bootstrap`: `false` (Add this setting **only** when initializing a fresh cluster with no data.)

Note: If the nodes in the cluster are identical in terms of disk layout, shared libraries, and so on, you can use the same copy of the `cassandra.yaml` file on all of them.

Example:

**`cluster_name: 'MyCassandraCluster'`**

**`num_tokens: 256`**

**`seed_provider:`**

**`- class_name: org.apache.cassandra.locator.SimpleSeedProvider`**

**`parameters:`**

**`- seeds: "10.168.66.41,10.176.170.59"`**

**`listen_address:`**

**`endpoint_snitch: GossipingPropertyFileSnitch`**

Note: Include at least one node from *each* datacenter.

5. In the `cassandra-rackdc.properties` file, assign the data center and rack names you determined in the Prerequisites. For example:

**Nodes 0 to 2**

# indicate the rack and dc for this node

**dc=DC1**

**rack=RAC1**

**Nodes 3 to 5**

# indicate the rack and dc for this node

**dc=DC2**

**rack=RAC1**

6. After you have installed and configured Cassandra on all nodes, start the seed nodes one at a time, and then start the rest of the nodes.

Note: If the node has restarted because of automatic restart, you must first stop the node and clear the data directories, as described [above](#).

Package installations:

**# sudo service cassandra start**

Tarball installations:

**# cd install\_location**

**# bin/cassandra**

7. To check that the ring is up and running, run:

Package installations:

**# nodetool status**

to create keyspace with particular replication in each datacenter

use “NetworkTopologyStrategy” instead of “SimpleStrategy” for create keyspace

```
> CREATE KEYSPACE "Excalibur" WITH REPLICATION = {'class' :  
'NetworkTopologyStrategy', 'dc1' : 3, 'dc2' : 2};
```

\*\*\*\*\*

**link ::**

<https://docs.datastax.com/en/cassandra/2.1/cassandra/initialize/initializeMultipleDS.html>

<https://www.packtpub.com/books/content/apache-cassandra-working-multiple-datacenter-environments>

commands for create keyspace in multi-datacenter ::

[https://docs.datastax.com/en/cql/3.1/cql/cql\\_reference/create\\_keyspace\\_r.html](https://docs.datastax.com/en/cql/3.1/cql/cql_reference/create_keyspace_r.html)

\*\*\*\*\*

\*\*\*\*\*

## *Cassandra Data partitioning*

\*\*\*\*\*

Cassandra is a distributed database that runs on multiple nodes. When you write data to the cluster, **partitioning scheme determines which node in the cluster stores that data**. For example, suppose you are inserting some data (Column-Value pair identified by a Row Key). Data partitioning protocol will dictate which node in the cluster is responsible for storing this data. Similarly, when you request data, the partitioning protocol will examine the Row Key and find the node in the cluster responsible for the row key and retrieve data from it.

Difference between Partitioning and Replication?

Data partitioning is concerned with picking a node in the cluster to store the first copy of data on. Replication determines number of additional nodes that will store the same data (for performance and fault tolerance reasons). Replication is discussed in the next section.

Partitioning => Picking out one node to store first copy of data on

Replication => Picking out additional nodes to store more copies of data

When you deploy a Cassandra cluster, you must assign a partitioner and assign each node an initial token value so each node is responsible for roughly an equal amount of data (load balancing). DataStax strongly recommends using the RandomPartitioner (default) for all cluster deployments.

To calculate the tokens for nodes in a single data center cluster, you divide the range by the total number of nodes in the cluster. In multiple data center deployments, you calculate the tokens such that each data center is individually load balanced. See [Generating Tokens](#) for the different approaches to generating tokens for nodes in single and multiple data center clusters.

Unlike almost every other configuration choice in Cassandra, the partitioner may not be changed without reloading all of your data. Therefore, it is important to choose and configure the correct partitioner before initializing your cluster. You set the partitioner in the [cassandra.yaml](#) file.

### *1) Murmur3Partitioner*

\*\*\*\*\*

*link ::*

[http://docs.datastax.com/en/cassandra/2.0/cassandra/architecture/architecturePartitionerM3P\\_c.html](http://docs.datastax.com/en/cassandra/2.0/cassandra/architecture/architecturePartitionerM3P_c.html)

\*\*\*\*\*

### *2) RandomPartitioner*

The RandomPartitioner uses tokens to help assign equal portions of data to each node and evenly distribute data from all the tables throughout the ring or other grouping, such as a keyspace. This is true even if the tables use different row keys, such as usernames or timestamps. Moreover, the read and write requests to the cluster are also evenly distributed and load balancing is simplified because each part of the hash range receives an equal number of rows on average. The RandomPartitioner (org.apache.cassandra.dht.RandomPartitioner) is the default partitioning strategy for a Cassandra cluster, and in almost all cases is the right choice. The RandomPartitioner distributes data evenly across the nodes using an MD5 hash value of the row key. The possible range of hash values is from 0 to  $2^{127} - 1$ .

\*\*\*\*\*  
link ::  
[http://docs.datastax.com/en/cassandra/2.0/cassandra/architecture/architecturePartitionerRandom\\_c.html](http://docs.datastax.com/en/cassandra/2.0/cassandra/architecture/architecturePartitionerRandom_c.html)  
[http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster\\_architecture/partitioning.html#data-distribution-in-the-ring](http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster_architecture/partitioning.html#data-distribution-in-the-ring)  
\*\*\*\*\*

## Cassandra Token Calculator

\*\*\*\*\*  
link ::  
<http://www.geroba.com/cassandra/cassandra-token-calculator/>  
\*\*\*\*\*

When creating a Cassandra cluster and not using virtual nodes that were introduced in version 1.2 (and are *not fully supported by OpsCenter yet*), you need to define the token range each individual cluster node is responsible for. If you set up Cassandra Clusters from time to time (like I do), this online calculator can be quite handy for you.

Just enter the amount of nodes you want to have initially and the partitioner (either RandomPartitioner or Murmur3Partitioner which is default in Cassandra 1.2+):

### Cassandra Token Calculator

Partitioner	<input type="radio"/> Murmur3Partitioner <input checked="" type="radio"/> RandomPartitioner
Number of nodes	<input type="text" value="2"/>
Result	0n85070591730234615865843651857942052864

[Calculate Tokens](#) Note: This Implementation uses the [BigInteger JavaScript Library](#) by Silent Matt.

### Explanation of token calculation

The token calculation is basically a function that divides the whole token range into equally sized subparts. For the RandomPartitioner, Cassandra offers a tool to calculate the partitions, for Murmur3, a Python script is provided:

- [Generating RandomPartitioner tokens](#)
- [Generating Murmur3Partitioner tokens](#)

**Recent Posts**

- ▶ This was Cassandra Meetup in Vienna
- ▶ Announcement: First Cassandra Meetup Vienna on January 13, 2014
- ▶ Cassandra Summit Europe 2013
- ▶ Cassandra Token Calculator
- ▶ How to Install Cassandra on Ubuntu 12.04 LTS in Windows Azure (with OpsCenter)

**Recent Comments**

- ▶ Zed on How to Install Cassandra on Ubuntu 12.04 LTS in Windows Azure (with OpsCenter)
- ▶ android on Cassandra Summit Europe 2013
- ▶ Datastax Cassandra repair service weird estimation and heavy load - CSS PHP on Cassandra Token Calculator
- ▶ This was Cassandra Meetup in Vienna - Geroba Data Technologies on Announcement: First Cassandra Meetup Vienna on January 13, 2014
- ▶ Rosie on Cassandra Token Calculator

here i created token for 2 node,  
0n85070591730234615865843651857942052864

*for server-node : 0*  
*for client-node : 85070591730234615865843651857942052864*

# configuration

The selection of partitioner need to be take before start the replication  
all other configuration parameters in this file same as i done earlier. Make changes only below lines  
in that configuration file.

## Step 1 - Calculate the token

```
*****
links ::
http://www.geroba.com/cassandra/cassandra-token-calculator/
*****
```

## Step 2 - Reconfigure the token properties in cassandra.yaml:

- comment out num\_tokens
- set the token assignment initial\_token
- leave auto\_bootstrap: true

*example ::*

```
# sudo service cassandra stop
```

```
# sudo nano /etc/cassandra/cassandra.yaml
```

```
num_tokens:
```

```
...
```

```
initial_token 85070591730234615865843651857942052864
```

```
....
```

```
auto_bootstrap: true
```

remove all data form server

```
# sudo rm -rf /var/lib/cassandra/data/*
```

```
# sudo mkdir /var/lib/cassandra/data/system
```

```
# sudo chown -R cassandra:cassandra /var/lib/cassandra/data
```

start the server and check the error log.

```
# service cassandra restart
```

```
# tailf /var/log/cassandra/system.log
```

```
# service cassandra status
```

if all working well, then :

(here i done with RandomPartitioner)



*# sudo nodetool status*

*Datacenter: dc1*

*=====*

*Status=Up/Down*

*|/ State=Normal/Leaving/Joining/Moving*

<i>--</i>	<i>Address</i>	<i>Load</i>	<i>Tokens</i>	<i>Owns (effective)</i>	<i>Host ID</i>	<i>Rack</i>
<i>UN</i>	<i>192.168.1.253</i>	<i>219.24 KiB</i>	<i>1</i>	<i>50.0%</i>	<i>21828972-91c2-403d-a228-1537b1630c99</i>	<i>rack1</i>
<i>UN</i>	<i>192.168.1.254</i>	<i>232.89 KiB</i>	<i>1</i>	<i>50.0%</i>	<i>f8fc523a-2786-4e74-8a82-75b0ffd67586</i>	<i>rack1</i>

*here the owns shows 50 %*

*\*\*\*\*\**

*link ::*

*detailed output info ::*

*<https://docs.datastax.com/en/cassandra/2.1/cassandra/tools/toolsStatus.html>*

*\*\*\*\*\**

*if there is any error while checking the log,*

*error :: [Unable to start Cassandra: "node already exists"](#)*

*# find / -name cassandra-env.sh*

*# nano /path/to/cassandra-env.sh*

*add this at bottom of file*

*JVM\_OPTS="\$JVM\_OPTS -Dcassandra.replace\_address=192.168.1.253"*

*# service cassandra restart*

*# tailf /var/log/cassandra/system.log*

*# service cassandra status*

*Don't forget to remove it once your done.*

*\*\*\*\*\**

*Link ::*

*<http://stackoverflow.com/questions/29323709/unable-to-start-cassandra-node-already-exists>*

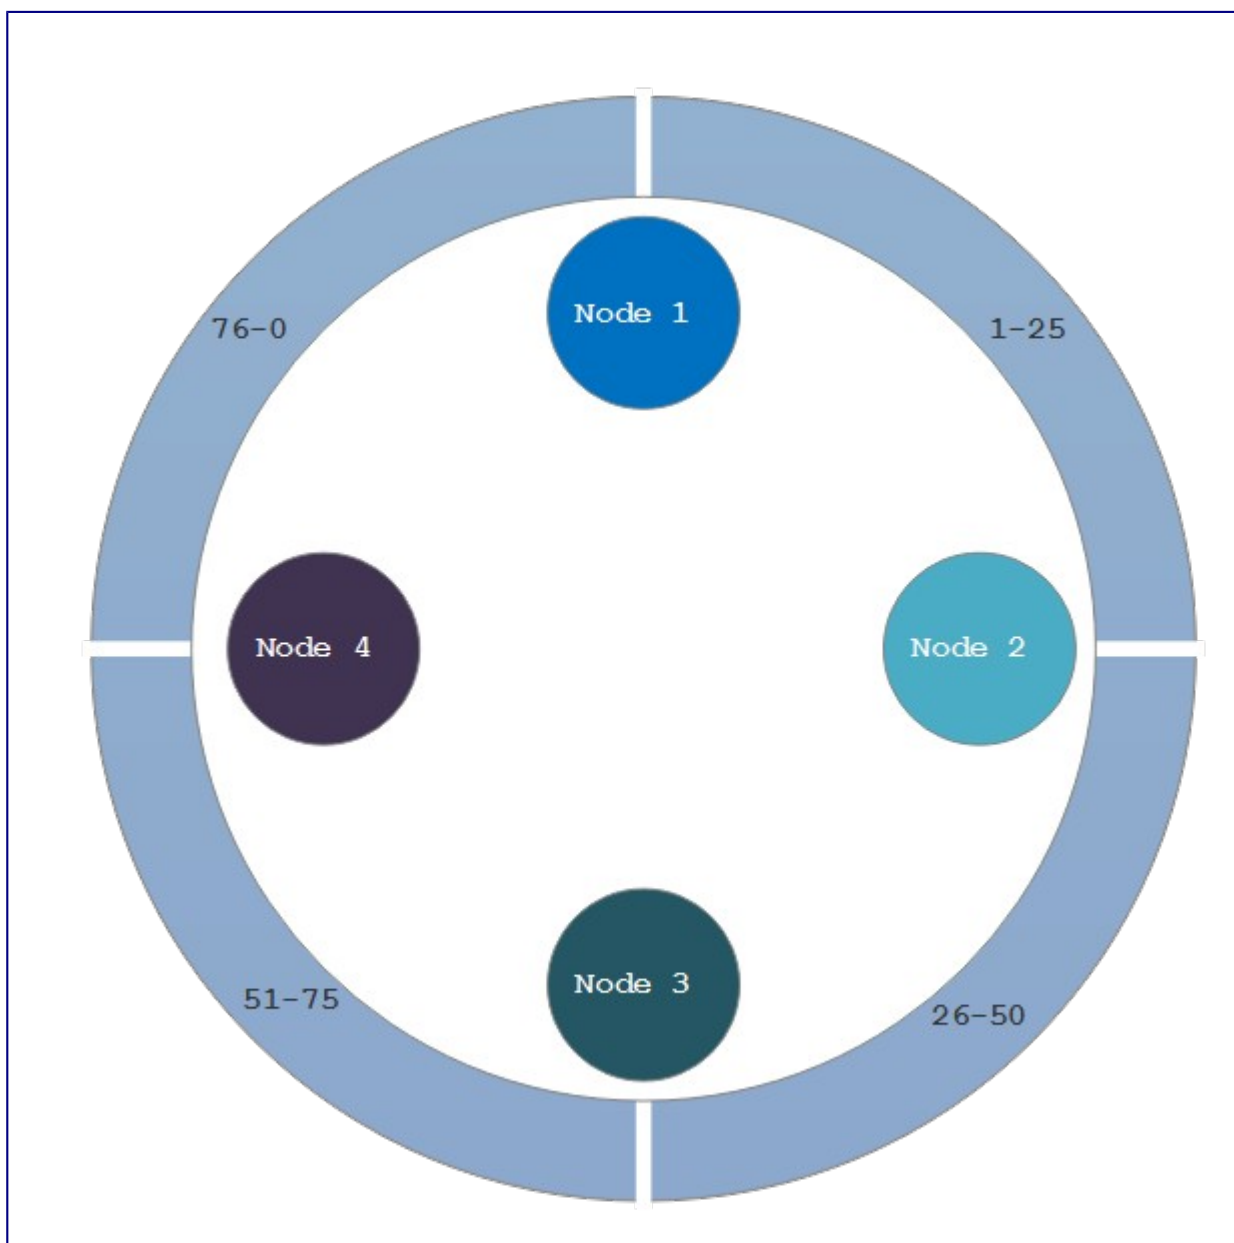
*\*\*\*\*\**

## ***Cassandra Partitioning & Clustering Keys Explained***

Primary Keys are defined when you create your table. The most basic primary key is a single column. A single column is great for when you know the value that you will be searching for

### **Partition Key**

The Partition Key is responsible for the distribution of data amongst the nodes. Let's look back to an earlier post on [Cassandra Data Model Basics](#), in which I described a four node cluster, as shown below. For simplicities' sake, let's assume hash values are between 0-100. When we insert the first row into the crossfit\_gyms table, the value of gym\_name will be hashed. Let's also assume that the first record will have a hash of 34. That will fall into the values that Node 2's partition is assigned. So the value of the Partition Key, 34, indicates the partition, 26-50, in the cluster/ring that the piece of data will be stored. Makes sense, huh?



## Compound Keys

But wait, there's more! Primary keys can also be more than one column. A multi-column primary key is called a Compound Key.

\*\*\*\*\*

## testing ::

\*\*\*\*\*

login to server

```
# cqlsh 192.168.1.253 9042
```

create a keyspace

```
> CREATE KEYSPACE biz WITH REPLICATION = { 'class' : 'SimpleStrategy',  
'replication_factor' : 1 };
```

check whether the keyspace created or not

```
> SELECT * FROM system_schema.keyspaces;
```

get into that keyspace

```
> use biz;
```

create a table with partition key for partitioning the data

```
> CREATE TABLE crossfit_gyms_by_location (  
    country_code text,  
    state_province text,  
    city text,  
    gym_name text,  
    PRIMARY KEY (country_code, state_province, city, gym_name)  
) WITH CLUSTERING ORDER BY (state_province DESC, city ASC, gym_name ASC);
```

add the values to the table for testing

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'CA', 'San Francisco', 'San Francisco CrossFit');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'CA', 'San Francisco', 'LaLanne Fitness CrossFit');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'NY', 'New York', 'CrossFit NYC');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'NY', 'New York', 'CrossFit Metropolis');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'NV', 'Las Vegas', 'CrossFit Las Vegas');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)  
VALUES ('USA', 'NV', 'Las Vegas', 'Kaizen CrossFit');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)
VALUES ('CAN', 'ON', 'Toronto', 'CrossFit Toronto');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)
VALUES ('CAN', 'ON', 'Toronto', 'CrossFit Leslieville');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)
VALUES ('CAN', 'BC', 'Vancouver', 'CrossFit Vancouver');
```

```
> INSERT INTO crossfit_gyms_by_location (country_code, state_province, city, gym_name)
VALUES ('CAN', 'BC', 'Vancouver', 'CrossFit BC');
```

check the data replicated on all node

```
> select * from crossfit_gyms_by_location;
```

country_code	state_province	city	gym_name
USA	NY	New York	CrossFit Metropolis
USA	NY	New York	CrossFit NYC
USA	NV	Las Vegas	CrossFit Las Vegas
USA	NV	Las Vegas	Kaizen CrossFit
USA	CA	San Francisco	LaLanne Fitness CrossFit
USA	CA	San Francisco	San Francisco CrossFit
CAN	ON	Toronto	CrossFit Leslieville
CAN	ON	Toronto	CrossFit Toronto
CAN	BC	Vancouver	CrossFit BC
CAN	BC	Vancouver	CrossFit Vancouver

(10 rows)

here i use “country\_code” as the partition key  
to see the hashed values of “country\_code”

```
> SELECT token(country_code),country_code FROM crossfit_gyms_by_location;
```

system.token(country_code)	country_code
11477211097623334992027430758387484140	USA
11477211097623334992027430758387484140	USA
11477211097623334992027430758387484140	USA
11477211097623334992027430758387484140	USA
11477211097623334992027430758387484140	USA
11477211097623334992027430758387484140	USA
145302877503818936703335776461659136135	CAN
145302877503818936703335776461659136135	CAN
145302877503818936703335776461659136135	CAN
145302877503818936703335776461659136135	CAN

```
> exit
```

*# nodetool getendpoints keyspace table key-value*

*ubuntu@ubuntu-client:/root\$ nodetool getendpoints biz crossfit\_gyms\_by\_location USA*  
*192.168.1.254*

*ubuntu@ubuntu-client:/root\$ nodetool getendpoints biz crossfit\_gyms\_by\_location CAN*  
*192.168.1.253*

*or search with the hashed value of key*

*ubuntu@ubuntu-client:/root\$ nodetool getendpoints biz crossfit\_gyms\_by\_location*  
*11477211097623334992027430758387484140*  
*192.168.1.254*

*ubuntu@ubuntu-client:/root\$ nodetool getendpoints biz crossfit\_gyms\_by\_location*  
*145302877503818936703335776461659136135*  
*192.168.1.253*

\*\*\*\*\*

**link ::**

**to add table ::**

<http://datascale.io/cassandra-partitioning-and-clustering-keys-explained/>

**data partitioning ::**

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster\\_architecture/partitioning.html#data-distribution-in-the-ring](http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster_architecture/partitioning.html#data-distribution-in-the-ring)

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/configuration/node\\_configuration.html#initial-token](http://docs.datastax.com/en/archived/cassandra/1.1/docs/configuration/node_configuration.html#initial-token)

**partitioning keys ::**

<https://jagadeeshs.wordpress.com/2015/07/15/cassandra/>

<https://dzone.com/articles/cassandra-data-modeling-primary-clustering-partiti>

**type of partitioner ::**

<https://10kloc.wordpress.com/2012/12/27/cassandra-chapter-4-data-partitioning/>

**token ::**

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token\\_generation.html#token-gen-cassandra](http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token_generation.html#token-gen-cassandra)

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token\\_generation.html#token-gen-cassandra](http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token_generation.html#token-gen-cassandra)

**partitioning configuration ::**

<https://support.datastax.com/hc/en-us/articles/205199715-Node-startup-fails-with-error-ConfigurationException-Cannot-change-the-number-of-tokens->

**visual testing ::**

<http://psanford.github.io/cassandra-visual-ring/>

**testing ::**

<http://stackoverflow.com/questions/30514237/what-node-does-cassandra-store-data-on>

\*\*\*\*\*

\*\*\*\*\*

## ***All links ::***

### ***tutorial ::***

[https://www.tutorialspoint.com/cassandra/cassandra\\_quick\\_guide.htm](https://www.tutorialspoint.com/cassandra/cassandra_quick_guide.htm)

### ***latest version check ::***

<http://cassandra.apache.org/download/>

### ***installation and configuration ::***

<https://www.digitalocean.com/community/tutorials/how-to-run-a-multi-node-cluster-database-with-cassandra-on-ubuntu-14-04>

<https://www.digitalocean.com/community/tutorials/how-to-install-cassandra-and-run-a-single-node-cluster-on-ubuntu-14-04>

### ***configuration directives ::***

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/configuration/node\\_configuration.html#initial-token](http://docs.datastax.com/en/archived/cassandra/1.1/docs/configuration/node_configuration.html#initial-token)

### ***keyspace creation ::***

[https://docs.datastax.com/en/cql/3.1/cql/cql\\_reference/create\\_keyspace\\_r.html](https://docs.datastax.com/en/cql/3.1/cql/cql_reference/create_keyspace_r.html)

<http://datascale.io/cassandra-partitioning-and-clustering-keys-explained/>

### ***multi-datacenter ::***

<https://docs.datastax.com/en/cassandra/2.1/cassandra/initialize/initializeMultipleDS.html>

<https://www.packtpub.com/books/content/apache-cassandra-working-multiple-datacenter-environments>

### ***commands for create keyspace in multi-datacenter ::***

[https://docs.datastax.com/en/cql/3.1/cql/cql\\_reference/create\\_keyspace\\_r.html](https://docs.datastax.com/en/cql/3.1/cql/cql_reference/create_keyspace_r.html)

### ***create keyspace for multi-data center ::***

[https://books.google.co.in/books?](https://books.google.co.in/books?id=aIro2eZEz4YC&pg=PT26&lpg=PT26&dq=Random+Partitioner+configuration+in+cassandra&source=bl&ots=qWhK0lXgOZ&sig=BZf3o_gM-laTIO0ayfd2z5Dsnvg&hl=en&sa=X&ved=0ahUKEwjkhMOO5sPQAhXDqY8KHRUSCh0Q6AEITzAH#v=onepage&q&f=false)

[id=aIro2eZEz4YC&pg=PT26&lpg=PT26&dq=Random+Partitioner+configuration+in+cassandra&source=bl&ots=qWhK0lXgOZ&sig=BZf3o\\_gM-laTIO0ayfd2z5Dsnvg&hl=en&sa=X&ved=0ahUKEwjkhMOO5sPQAhXDqY8KHRUSCh0Q6AEITzAH#v=onepage&q&f=false](https://books.google.co.in/books?id=aIro2eZEz4YC&pg=PT26&lpg=PT26&dq=Random+Partitioner+configuration+in+cassandra&source=bl&ots=qWhK0lXgOZ&sig=BZf3o_gM-laTIO0ayfd2z5Dsnvg&hl=en&sa=X&ved=0ahUKEwjkhMOO5sPQAhXDqY8KHRUSCh0Q6AEITzAH#v=onepage&q&f=false)

### ***partitioning ::***

<https://10kloc.wordpress.com/2012/12/27/cassandra-chapter-4-data-partitioning/>

<https://support.datastax.com/hc/en-us/articles/205199715-Node-startup-fails-with-error-ConfigurationException-Cannot-change-the-number-of-tokens->

<http://distributeddatastore.blogspot.in/2015/07/cassandra-data-partitioning-using.html>

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster\\_architecture/partitioning.html#data-distribution-in-the-ring](http://docs.datastax.com/en/archived/cassandra/1.1/docs/cluster_architecture/partitioning.html#data-distribution-in-the-ring)

***token generation ::***

<http://www.geroba.com/cassandra/cassandra-token-calculator/>

[http://docs.datastax.com/en/cassandra/2.0/cassandra/configuration/configGenTokens\\_c.html](http://docs.datastax.com/en/cassandra/2.0/cassandra/configuration/configGenTokens_c.html)

[http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token\\_generation.html#token-gen-cassandra](http://docs.datastax.com/en/archived/cassandra/1.1/docs/initialize/token_generation.html#token-gen-cassandra)

***partion key and clustering key ::***

<https://jagadeeshs.wordpress.com/2015/07/15/cassandra/>

<https://dzone.com/articles/cassandra-data-modeling-primary-clustering-partiti>

***error solution for already existing data if initial\_token changed ::***

<http://stackoverflow.com/questions/29323709/unable-to-start-cassandra-node-already-exists>

***detailed infoemation in partitioning ::***

<https://docs.datastax.com/en/cassandra/2.1/cassandra/tools/toolsStatus.html>

***partitioning testing ::***

<http://stackoverflow.com/questions/30514237/what-node-does-cassandra-store-data-on>

***visual testing of nodes ::***

<http://psanford.github.io/cassandra-visual-ring/>

\*\*\*\*\*